

*Escribe aquí
tu frase favorita.*

E indica aquí su autor

Agradecimientos

Me gustaría agradecer...

También quiero destacar...

Por último...

Índice general

Índice de figuras	III
Índice de tablas	IV
Índice de algoritmos	V
Resumen	1
1. Introduction	3
1.1. Context	3
1.2. Objectives	3
1.2.1. Main objective	3
1.2.2. Specific objectives	3
2. Theoretical framework	5
2.1. Problem Definition	5
2.2. Mono and Multi-modal Image Registration	5
2.3. Rigid and Non-rigid Registration	5
2.4. Image Warping	6
2.5. Diffeomorphism	6
2.6. Classical Methods	6
2.7. Deep learning-based methods	6
2.7.1. U-Net	6
2.8. Transformers	6
2.8.1. Mechanism of Self-attention	6
2.8.2. Encoder Block	7
2.8.3. Decoder Block	7
2.8.4. Multi-head Self-attention	7
2.9. Vision Transformers	7
2.9.1. SymTrans	8
2.9.2. TransMorph	8



3. Methodology	11
3.1. Proposed Architecture	11
3.2. Image Similarity Metric	11
4. Experiments	13
4.1. Dataset and Metrics	13
4.2. Baseline Methods	13
4.3. Implementation Details	13
4.4. Results	13
4.4.1. Quantitative Results	13
4.4.2. Qualitative Results	13
4.4.3. Computational Complexity	13
5. Conclusion and Limitations	14
5.1. Conclusion	14
5.2. Limitations	14
A. Annex A	16
Bibliografía	17

Índice de figuras

2.1. Diagram of Vision Transformer. Adapted from 1	8
2.2. Visualization of Self-Attention on Input Images. Reproduced from 1	9



Índice de tablas

Índice de algoritmos



Abstract

Introduction

1

Escribe aquí la introducción de tu Trabajo Fin de Máster, utilizando tantas secciones, subsecciones y subsubsecciones como estimes necesarias.

1.1. Context

Esta palabra está en negrita. *Esta palabra* está en cursiva. **Esta palabra** se destaca en púrpura.

1.2. Objectives

En la sección ?? se muestran ejemplos de palabras en negrita, cursiva y destacadas en púrpura.

Una Red Generativa Antagónica o *Generative Adversarial Network* (GAN) es... (?).

? diseñaron las redes generativas antagónicas como...

1.2.1. Main objective

1.2.2. Specific objectives

1. Objetivo parcial 1.

2. Objetivo parcial 2.

3. Objetivo parcial 3.

Theoretical framework

2.1. Problem Definition

The problem of image registration is an optimization problem, and its candidate solutions are all the possible transformations that maximize or, equivalently, minimize a chosen (di)similarity measure that quantifies the image similarity between the fixed and moving image; poses as the target function.

Depending on the particular problem setting (e.g. deformable registration), we will either face a linear or nonlinear optimization problem, which will affect the possible solution strategies.

2.2. Mono and Multi-modal Image Registration

2.3. Rigid and Non-rigid Registration

The rigid registration setting, the resulting transformation is expressed as a matrix mapping each point of the moving image to the target / fixed image. The possible transformations expressed by the matrix are rigid and affine transformations. Rigid transformations are limited to rotations and translations, and affine transformations extend the rigid setting, allowing for stretching and skewing.

In the non-rigid case, limitations are lifted, and allows for arbitrary mappings of individual image pixels, producing displacement fields.

The problem with displacement fields is that, depending on the setting, only a subset of all possible transformations are indeed plausible, so the definition is extended to accommodate for this. A particular deformation is said to be plausible if the pixel movement is smooth, simulating natural elastic deformation.

Usually, regularization techniques are employed in order to enforce the aforementioned properties.

2.4. Image Warping

2.5. Diffeomorphism

2.6. Classical Methods

Demons, other iterative methods...

2.7. Deep learning-based methods

When approaching the original problem of image registration... When using deep learning, we don't directly calculate the

Differences between unsupervised and supervised deep learning methods.

2.7.1. U-Net

Some paper of implementation using U-Net or similar model

2.8. Transformers

Transformers were introduced ...

Currently, state of the art in sequence processing, tasks such as NLP...

Encoder vs Decoder.

2.8.1. Mechanism of Self-attention

When using transformers, the input sequence is tokenized, obtaining a set of tokens - a collection of distinct, unordered elements.

The next step is the projection of the tokens into a distributed geometrical space of continuous-valued vectors - what is referred to as an embedding. This is done in order to preserve the semantical relationships among the tokens.

In this low-dimensional space, we expect that the projections of the original tokens which hold stronger semantic relationships with each other are, indeed, closer to each other than their counterparts.

However, after the tokenization step, we have effectively lost the order of the original sequence. In order to maintain the notion of order necessary to process the input as a sequence, transformer blocks use positional encoding.

Positional encoding alters the embeddings depending of the position of the token in the original sequence.

One of the possible techniques for implementing positional encoding is using the sinusoidal function.

We then perform feature-based attention on the resulting embeddings. The inductive bias by which feature-based attention is based upon, is to allow the network to place its attention not only based in the original order of the sequence, but taking into account the content of the tokens.

In order to quantify the similarity of tokens, we will calculate the vector similarity among the tokens' projections into the learned embedding space. As this is a low-dimensional space, we can make use of the inner vector product.

The transformer uses 3 different representations of the embedding matrix, the queries, keys and values.

For each token (i.e. vector), we create a query vector, key vector and value vector. The query, key and value matrices are learned as part of the training process. In order to reduce the computational complexity of the operation, this resulting vectors are usually in lower-dimensional spaces.

The concepts of key, value and query are originally part of information retrieval systems.

Self attention is performed for each position.

The main idea is that we compute the similarity of each token of the sequence with each other token. Then, the result of the self-attention is a vector, result of the sum of each token in the sequence multiplied by its similarity score. In this way, we are capturing the most important context individually for each of the tokens in the sequence.

Self-attention with matrices First, we obtain the query, key and value matrices by multiplying the embedding matrix with the learned weight matrices W_Q , W_K , W_V . Each of the rows of the embedding matrix corresponds to a token of the input sequence.

$$\text{softmax}\left(\frac{QxK^T}{\sqrt{d_k}}\right)V = Z$$

The softmax outputs a probability distribution, with all of its components adding to one.

This concept of self-attention is further developed with multi-headed attention.

2.8.2. Encoder Block

2.8.3. Decoder Block

2.8.4. Multi-head Self-attention

With multi-headed attention, we repeat the original self-attention process, obtaining multiple representational spaces, in the form of multiple sets of query, key and value matrices. Theoretically, this allows the model to perform attention in different, independent low-dimensional spaces, which translates to the ability to jointly perform attention from different representation subspaces, at different positions.

2.9. Vision Transformers

In [Dosovitskiy et al. \(2021\)](#), vision transformers are introduced.

- Parallelizable - Bigger receptive fields compared to CNN

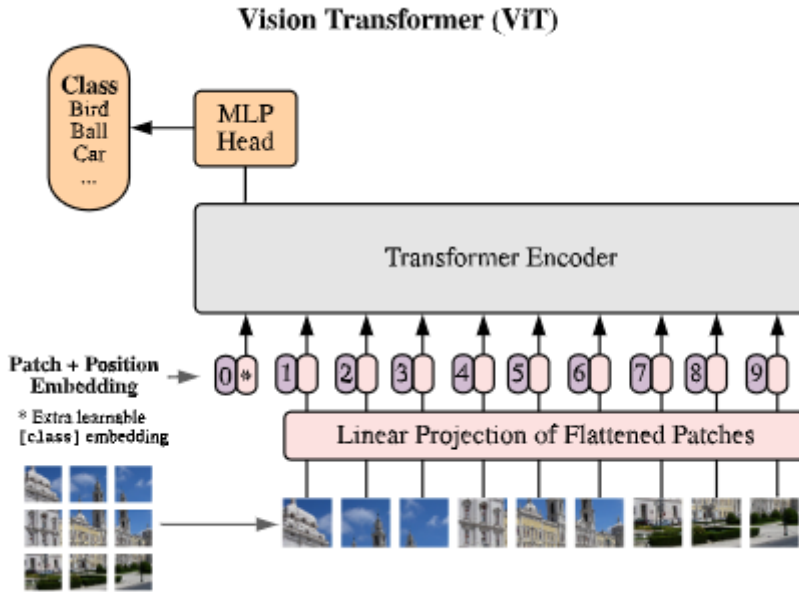


Figure 2.1: Diagram of Vision Transformer. Adapted from [1](#)

In order for a transformer network to be able to handle 2D images, they propose a method to convert 2D images into sequences by splitting the original image into smaller 2D patches, which are then further processed as described in [2.8.2](#).

Formally, the original image $x \in \mathbb{R}^{H \times W \times C}$ is reshaped into a sequence of flattened 2D patches $x_p \in \mathbb{R}^{N \times P^2 \times C}$, where (H, W) is the original resolution of the image, C denotes the number of channels, (P, P) is the patch resolution, and $N = HW/P^2$ is the total patch number (i.e. sequence length).

Also, the paper explores the possibility of feature maps obtained from a CNN as the input sequence, as an alternative to the aforementioned image patches.

In terms of the positional embeddings, they employ 1D positional embeddings. They also tried using 2D positional encodings without showing any further improvements in the model's performance.

One of the most important discoveries of the paper is that, upon inspection, the attention distance (i.e. average distance in image space across which information is integrated), analogous to the receptive field size of CNNs, is far greater than in its counterpart.

The results empirically confirm that the ability of performing of global attention, enabled by the self-attention mechanism, is used by the model in practice, where some of the attention heads attend most of the image already in the lower layers.

2.9.1. SymTrans

2.9.2. TransMorph




Figura 2.2: *Visualization of Self-Attention on Input Images. Reproduced from [1](#)*

Methodology

3

3.1. Proposed Architecture

3.2. Image Similarity Metric

Experiments

4

4.1. Dataset and Metrics

4.2. Baseline Methods

4.3. Implementation Details

4.4. Results

4.4.1. Quantitative Results

4.4.2. Qualitative Results

4.4.3. Computational Complexity

Conclusion and Limitations

5

5.1. Conclusion

5.2. Limitations

Annex A



Bibliografía

Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., y Houlsby, N. (2021). An image is worth 16x16 words: Transformers for image recognition at scale.